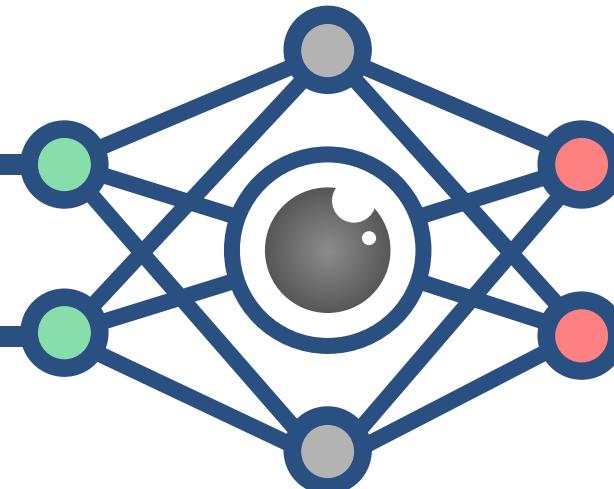


CS3485

Deep Learning for Computer Vision




Lec 8: Data Augmentation and Deep CNNs

Announcement


- Lab2:
 - Lab grades are out, let me know if you have any questions!
 - Don't forget the default parameters!
- Quiz #3 later today!

(Tentative) Lecture Roadmap

Basics of Deep Learning



Deep Learning and Computer Vision in Practice

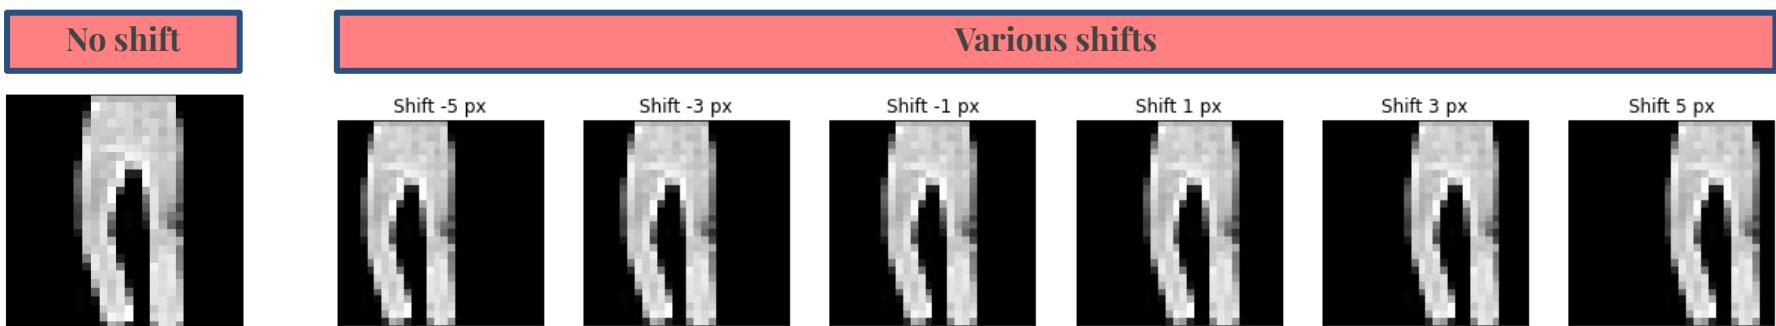


Ways to improve

- Last time, we saw that we can improve the classification task in the FashionMNIST dataset by using **Convolutional Neural Networks**.
- Despite our final classification outcome being pretty good, we can still improve it in some ways that we haven't tried last time:
 - By adding **regularization** (dropout, for example) and **Batch Normalization** to the network.
 - By training the network for **longer** (more than 5 epochs).
 - By tuning some of the network constants (also called **hyperparameters**), such as the optimizer's learning rate, the batch size, the number of strides and the padding of each ConvLayer.
 - By trying different amount of units/filters per layer to be learned.
 - By using **data augmentation**.
 - By adding **more layers** and making the network able to learn more complex image features.
- Today, we'll focus our efforts on the last two options: we'll see how making the **data (the input)** or the **network (the model)** "richer" can improve our classification performance.

Issues with Shifting

- Last time, we saw that CNNs do much better at classifying Fashion MNIST data than simple Multilayer Perceptrons.
- Today, we'll check how well their classifier works when we slightly change some of the images in a way that their classes would still be recognizable.
- This happens when you shift the image below some pixels to the right and to the left:



In these examples, the original class (“trouser”) shouldn't become less recognizable because of the shifts.

Trying out the CNN on the shifted images

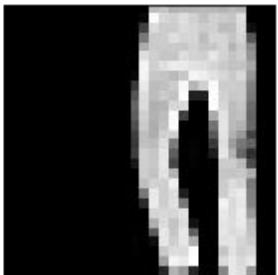
- Let's see how the model trained in the last class predicts the classes of the trouser shifted **1 to 5 pixels to the right and to the left**:

```
softmax = nn.Softmax() # Define the softmax function (remember that the more does not output probs.)  
idx = 24300 # The index of the trouser from the last slide  
  
preds = []  
for ix in range(-5, 6):  
    img = x_train[ix] # Read the desired image  
  
    img_rolled = np.roll(img, px, axis=1) # Roll the image by "px" pixels  
  
    img_rolled = torch.Tensor(img_rolled / 255.) \  
                 .view(-1, 1, 28, 28) \  
                 .to(device)           # Scale and reshape the image to the image  
                           # format used during learning. Register the  
                           # result to the GPU  
  
    pred = softmax(model(img_rolled)) # Apply our learned model to predict the class probabilities  
  
    preds.append(pred.cpu().detach().numpy()) # Post process the prediction and then save it to the list
```

* In the code above, we are using the variable names and libraries from the previous class. It's like its continuation.


Trying out the CNN on the shifted images

- Now we can plot the probabilities of each shifted image to belong to each of the **10** possible classes.
 - For most shifts, the network finds the right class “trouser”.
 - But, unexpectedly, the network **makes very bad guesses** for the images shifted closer to the border of the image.




- In fact, **it seems to be sure** that the image on the left is a *sandal*!
 - What can we do to fix this?

Prob. of each class for various shifts (CNN)




(A digression) CNNs and their receptive field

- To be fair with the CNN model, it does **quite a good work** when compared to the Multilayer Perceptron model (on the right).
- The improvement CNN adds to the pure fully connected MLP is related to the **receptive field** of convolutional and pooling layers.



- This means that later individual units have information about greater areas of the original image.
- This enables capturing of some shifting.



Data augmentation as a solution

- **But back to CNNs!** We noticed that these issues with image shifting can have on a model's prediction accuracy.
- However, in the real world, we might encounter various scenarios, such as the following:
 - Images are rotated slightly,
 - Images are zoomed in/out (scaled),
 - Some amount of noise is present in the image,
 - Images have low brightness,
 - Images have been flipped,
 - Images have been sheared (one side of the image is more twisted).
- A neural network that does not take the preceding scenarios into consideration won't provide accurate results.
- One solution to that issue is to **artificially change the data** in the dataset in a way to consider the above settings. This is called **Data Augmentation***.

* In other contexts, augmentation can also mean “make the dataset larger”, but in the end of the day, it is the same as we are doing.

Data augmentation via transformations

- The strategy we'll take consists in making random changes in each of our datapoints before they enter in our `train_batch` function.
- We'll use the very handy `transforms` from `torchvision`, usually imported as:

```
from torchvision.transforms import transforms
```

- A useful tool found in there is the **affine transformation** using `RandomAffine`:

```
transforms.RandomAffine(degrees, translate=None, scale=None, shear=None)
```

whose objects are python functions (`nn.Modules`, in fact) that perform either a random rotation, translation, scaling, shearing or any subset of them. Its parameters are:

- `degrees` (a number): Range of degrees to select from
- `translate` (a tuple): Maximum image fraction for horizontal and vertical shifts.
- `scale` (a number or a tuple): Scaling factor interval
- `shear` (a number): Range of pixels in the image will be sheared horizontally.

Examples of affine transformations

- Here are some examples of random affine transformations on an image from Fashion MNIST using `transforms.RandomAffine`*:



*Check the documentation [here](#) for more details on the layer and on other possible parameters.

Other Transformations

- The `transforms` library also provides more options of transformations*. For example:

- Change the perspective (`transforms.RandomPerspective`):



- Cropping a part of the image out (`transforms.RandomCrop`):



- Add Gaussian noise (`transforms.GaussianBlur`):



*[Here](#) you can find a list of all possible transforms available in PyTorch.

Other Transformations

- Invert the grayscale values / colors (`transforms.RandomInvert`):



- We can compose many different transformations using `transforms.Compose` that receives a list of `transforms` modules and processes them sequentially on the data.
- For example, the following code generates a transformation that first randomly rotates an image and then randomly inverts its colors.

```
transforms.Compose([transforms.RandomAffine(180), transforms.RandomInvert()])
```



Adding a transformation to the dataset


- The simplest way to add a transformation to the dataset is to apply it in the `__getitem__` function to the image being gotten.
- This way, this random transformation will happen whenever the DataLoader is fetching the data to compose the mini-batch.
- In our example, we wish the network to learn that horizontal shifts shouldn't change the object's class.
- Therefore we can augment the dataset by applying random horizontal shifts to the images.

```
from torchvision.transforms import transforms

class FMNISTDataset(Dataset):
    def __init__(self, x, y):
        x = x.view(-1, 1, 28, 28).float()/255
        self.x, self.y = x, y
        self.shift = transforms.RandomAffine(0, translate=(0.5, 0))
    def __getitem__(self, ix):
        x = self.x[ix]
        x = self.shift(x)
        return x.to(device), self.y[ix].to(device)
    def __len__(self):
        return len(self.x)
```

Result of the augmentation

- By making just that change, we are able to achieve the result for the same trouser image from before.
 - Notice that the network became more “invariant” to horizontal shifts, as it makes the right prediction with certitude despite the shifts.
 - This, however, came at a price:
 - a. Adding the random shifting operation at each `getitem` made the overall 5 epoch learning take 6 min (from 53s).
 - b. The new test accuracy is at around 88% (from 91% from before)



Problems with augmentation

- The problem **a** is easy to fix, as the purpose of the previous code was only to serve as an illustration of the augmentation process.
- In PyTorch there are ways to make the transformation application more efficient, by, for example, using them right when you load the data.

```
transform_train = transforms.Compose([transforms.RandomAffine(0, translate=(0.5, 0)),
                                     transforms.ToTensor()])
fmnist_train = datasets.FashionMNIST('~/FMNIST', download=True, train=True, transform=transform_train)
```


and them changing other parts of the code so we don't need to instantiate our own `Dataset` object, which is inefficient (*these details go beyond the scope of our course*).

- Problem **b**, however, is harder to solve, since an augmented dataset **is intrinsically richer and more complex** than the original data.
- Typically, it'd require at least **going through more training epochs** or **changing the network** to more complex ones.


Exercise (*In pairs*)

- For each image classification problem below, list which augmentations make sense and which are risky (not ideal) or unnecessary:


Satellite Images (for what type of region)



Handwritten Digits



Traffic Sign Classification




Cell (various types of white blood cell)



Making the model more complex

- We just saw that it is possible to learn a better classification model by presenting a richer variety of data, **even if that data is artificially augmented**.
- Another way to come up with a better model is by training a network whose feature learning phase can capture **more nuanced and representative visual features**.
- With such these more complex features, we hope that the final densely connected layers will be able to output good classifications.



Making the network deeper

- How to come up with better feature learners?
- Over the recent years, researchers have noticed that simply adding more ConvLayers before the dense classifier usually bring improvements.
- This pursuit of more layered nets gave rise to what is known as **Deep Learning**, which is, simply put, **the feature learning process that uses multilayered neural networks**.
- In other words, deep learning is, in many ways, just **representation learning**
- *Later in the course, we'll see why going deeper helps learning.*



The ImageNet Dataset


- Historically, Deep Learning started to impress the world in 2012, when a deep net called AlexNet broke the classification record on the ImageNet dataset.
- This dataset spans *1000* classes and contains *1,281,167* training and *100,000* test images* of various sizes.
- The **images are very realistic**, all hierarchically annotated by humans.



*In fact, this is just a subset of +14 million images spanning more than 20k classes called the ImageNet project. More info on it [here](#).

The ImageNet Challenge

- Since 2012, Deep Learning has outperformed every other method in the ImageNet's Top 5* Classification competition.
- Starting from 2014, it also **overcame humans**** when submitted to the same challenge.
- One common feature of all these winning networks is that they were **getting deeper and deeper**.
- Today we'll focus on one of the runner-ups from the 2014 edition: the **VGG16 network**.




* The true class only need to be among the top 5 predicted classes to be considered a successful prediction.

** Note that these methods need to identify 1 of a 1000 possible classes, while humans can recognize a much larger number of categories.


The VGG16 Network

- The VGG16 net, for Visual Geometry Group (VGG) at University of Oxford, who developed the network in 2014, is a simple, by very deep network, with 16 layers!
- While the input RGB image has to be reshaped to 224×224 pixels, it uses many ConvLayers and max-poolings to gradually decrease its size, before the dense layers.



VGG16 in PyTorch

- In a simplified way, the VGG16 can be summarized as follows:



- Although I'm sure you can code that network up from scratch, PyTorch also provides the model as it was conceived via in `touchnvision`:

```
from torchvision import models  
model = models.vgg16()
```

The summary of VGG16

```
from torchsummary import summary  
summary(model.to(device), (3, 224, 224))
```

Layer (type)	Output Shape	Param #
<hr/>		
Conv2d-1	[-, 64, 224, 224]	1,792
ReLU-2	[-, 64, 224, 224]	0
Conv2d-3	[-, 64, 224, 224]	36,928
ReLU-4	[-, 64, 224, 224]	0
MaxPool2d-5	[-, 64, 112, 112]	0
Conv2d-6	[-, 128, 112, 112]	73,856
ReLU-7	[-, 128, 112, 112]	0
Conv2d-8	[-, 128, 112, 112]	147,584
ReLU-9	[-, 128, 112, 112]	0
MaxPool2d-10	[-, 128, 56, 56]	0
Conv2d-11	[-, 256, 56, 56]	295,168
ReLU-12	[-, 256, 56, 56]	0
Conv2d-13	[-, 256, 56, 56]	590,080
ReLU-14	[-, 256, 56, 56]	0
Conv2d-15	[-, 256, 56, 56]	590,080
ReLU-16	[-, 256, 56, 56]	0
MaxPool2d-17	[-, 256, 28, 28]	0
Conv2d-18	[-, 512, 28, 28]	1,180,160
ReLU-19	[-, 512, 28, 28]	0

A new type of layer.

Conv2d-20	[-, 512, 28, 28]	2,359,808
ReLU-21	[-, 512, 28, 28]	0
Conv2d-22	[-, 512, 28, 28]	2,359,808
ReLU-23	[-, 512, 28, 28]	0
MaxPool2d-24	[-, 512, 14, 14]	0
Conv2d-25	[-, 512, 14, 14]	2,359,808
ReLU-26	[-, 512, 14, 14]	0
Conv2d-27	[-, 512, 14, 14]	2,359,808
ReLU-28	[-, 512, 14, 14]	0
Conv2d-29	[-, 512, 14, 14]	2,359,808
ReLU-30	[-, 512, 14, 14]	0
MaxPool2d-31	[-, 512, 7, 7]	0
AdaptiveAvgPool2d-32	[-, 512, 7, 7]	0
Linear-33	[-, 4096]	102,764,544
ReLU-34	[-, 4096]	0
Dropout-35	[-, 4096]	0
Linear-36	[-, 4096]	16,781,312
ReLU-37	[-, 4096]	0
Dropout-38	[-, 4096]	0
Linear-39	[-, 1000]	4,097,000
<hr/>		
Total params: 138,357,544		
Trainable params: 138,357,544		
Non-trainable params: 0		
<hr/>		
(...)		

Adaptative Average Pooling and Other VGG's

- As you may have noticed on the previous summary*, VGG16 utilizes a layer we haven't yet learned, the **Adaptive Average Pooling layer**.
- It is similar to `nn.AvgPool2d`, which returns the average of a section instead of the maximum, which `nn.MaxPool2d` does. In both cases, we set choose the kernel size.
- In `nn.AdaptiveAvgPool2d`, we instead set the output size, and it automatically computes the kernel size so that the specified size is returned.
- This layer plays an important role **in the transition from the feature learning phase to the classifier** and will be important in our next class.
- This layer is found in other models, such as VGG16's "siblings": VGG13 and VGG19, width 13 and 19 layers, respectively, which can be used via `models.vgg13()`, and `models.vgg19()`.

(...)	(...)	(...)
ReLU-30	[-1, 512, 14, 14]	0
MaxPool2d-31	[-1, 512, 7, 7]	0
AdaptiveAvgPool2d-32	[-1, 512, 7, 7]	0
Linear-33	[-1, 4096]	102,764,544
(...)	(...)	(...)

* Despite not explicitly showing here, there is a flattening layer in between the AvgPool and the Linear layers, as its official [implementation](#) recognizes.

The challenges of Deep Nets

- Note that in VGG16 we have to train **more than 135 million parameters** on **RGB images of size 224×224!**
- Using a simple GPU, we were taking ~1 min to learn *800k* weights for just *5* epochs on *60000* grayscale images of size *28×28*.
- For most applications, **it is not worth** to retrain these networks, especially if one is running on a low computational/memory budget.
- Also, the dataset VGG16 was trained on (ImageNet) has *+1* million images to be trained on.
- Two issues that are very common in most deep learning applications:
 - a. The **models are huge** and most companies can't afford the of computational requirement.
 - b. These models need to be **trained on very large datasets** so to justify their complexity. In many applications, the datasets are very small (*one could recur to data augmentation in this case*).
- *Next class*, we'll see how we can still leverage the capacities of deep learning models in the applications at a considerably low computational cost.

Exercise (*In pairs*)

- Go back the VGG16's [summary](#) and explain how the output sizes change as they do (remember that each ConvLayer uses 3×3 kernels). *Hint:* try to **print** the model and see if it gives you any help:

```
from torchvision import models  
model = models.vgg16()  
print(model)
```